

Ontology Engineering: Current State, Challenges, and Future Directions

Tania Tudorache

Stanford Center for Biomedical Informatics Research, Stanford University, CA, USA

E-mail: tudorache@stanford.edu

Abstract.

In the past decade, we have witnessed a significant adoption of ontologies in a variety of fields including biomedicine, finance, engineering, law, and cultural heritage. The ontology engineering field has been strengthened by the adoption of several standards pertaining to ontologies, by the development or extension of ontology building tools, and by a wider recognition of the importance of standardized vocabularies and formalized semantics. Research into ontology engineering has also produced methods and tools that are used more and more in production settings. Despite all these advancements, ontology engineering is still a difficult process, and many challenges still remain to be solved. This paper gives an overview of how the ontology engineering field has evolved in the last decade and discusses some of the unsolved issues and opportunities for future research.

Keywords: Ontologies, Ontology Engineering, Methods, Standards, Tooling, Patterns, Challenges, Future Research

1. Ontologies Make an Impact

The research on ontologies in computer science started in the early 1990s. Ontologies were proposed as a way to enable people and software agents to seamlessly share information about a domain of interest. An ontology was defined as a conceptual representation of the entities, their properties and relationships in a domain [1]. The ultimate goal of using ontologies was to make the knowledge in a domain computationally useful [2]. The initial research period was followed by a time of great excitement about using ontologies to solve a wide range of problems. However, the enthusiasm dwindled in the early 2000s, as the methods and infrastructures for building and using ontologies were not mature enough at that time. Nonetheless, significant changes have taken place in the last decade: The research and development on ontologies had a big boost, more standardization efforts were on the way, and industry started to buy into semantic technologies. As a result, ontologies are now much more widely adopted in academia, industry and government environments, and are finally making an impact in many domains.

Biomedicine has widely adopted ontologies since their beginnings. The Gene Ontology (GO) [3]—a comprehensive ontology describing the function of genes—is the poster child for a successful ontology development project that has produced a big impact in biomedical research. Indeed, GO is routinely used in the computational analysis of large-scale molecular biology and genetics experiments [4]. Researchers have also used ontologies in biomedicine to standardize terminology in particular domains, to annotate large biomedical datasets, to integrate data, and to aid structured data mining and machine learning [5, 6].

One notable example of the impact ontologies are making in healthcare is the development of the 11th revision of the International Classification of Diseases (ICD-11). ICD—developed by the World Health Organization (WHO)—is the international standard for reporting diseases and health conditions, and is used to identify health trends and statistics on a global scale [7]. ICD-11 is now using OWL to encode the formal representation of diseases, their properties, and relations, as well as mappings to other terminologies [8].

The **financial industry** has embraced the use of ontologies. The most prominent example is the Financial Industry Business Ontology (FIBO)—the industry

standard resource for the definitions of business concepts in the financial services industry [9]. FIBO is developed by the Enterprise Data Management Council (EDMC) and it is standardized through the Object Management Group (OMG). FIBO is built as a series of OWL ontologies and is developed using a rigorous and well-defined process, known as the “Build-Test-Deploy-Maintain” methodology.

Engineering is another field that has adopted ontologies from the early 1990s, long before the standardization of the current Semantic Web languages, such as OWL and RDF [10, 11]. In the last decade, we have witnessed significant efforts around using ontologies to cover different aspects of engineering ranging from defining requirements [12], to integrating different engineering models [13], to detecting inconsistencies in models in multidisciplinary engineering projects [14]. Sabou et. al [15] provide a comprehensive overview into how ontologies and Semantic Web technologies can assist in building intelligent engineering applications.

A lot of work has gone into developing methods and tools for publishing linked datasets of the vast **cultural heritage field** in the last decade [16]. One outstanding example is the linked open dataset of the Europeana Collections¹ [17] which provides access to millions of artworks, artefacts, books, films and music from European museums, galleries, libraries and archives. Another example comes from the scholarly publishing domain, in which the SPAR (Semantic Publishing and Referencing) Ontologies [18] are having a high impact since they were released in 2015. The recently released Italian Cultural Heritage knowledge graph, ArCO² [19], consists of a network of seven high-quality ontologies, modeling the cultural heritage domain, and contains over 169 million triples about 820 thousand cultural entities. The testament to the importance of ontologies in the cultural heritage field is shown also by the adoption of the ISO 21127:2014 [20] standard that prescribes an ontology allowing the exchange of cultural heritage data between institutions.

Other fields have also adopted ontologies more widely. Researchers have used ontologies in the **legal domain** to formally represent laws and regulations, to simulate legal actions, or for semantic searching and indexing [21]. In the **geographical domain**, the ISO

19150-1:2012 [22] defines a high-level model of the components required to handle semantics in the ISO geographic information standards with ontologies. The second part of the standard, ISO 19150-2:2015 [23] defines rules to convert the UML models used in the ISO geographic information standards into OWL.

The examples we mentioned above are not meant to be comprehensive. They show how ontologies have been embraced by a wide range of fields in the last decade and how they are making an impact.

This paper is meant to give a retrospective overview of how the ontology landscape and ontology engineering have evolved in the last decade, current challenges, and prospects for future research. This paper can hopefully also serve as an introduction for newcomers in the field. We briefly discuss standards relevant to ontology engineering that have been adopted in the last decade (Section 2), highly visible and influential ontologies and knowledge bases that are constructed by large communities (Section 3), trends in ontology engineering from the last ten years (Section 4), and current challenges (Section 5) and opportunities for future research (Section 6).

2. New Standards

The significant standardization efforts on ontologies and Semantic Web languages in the last decade also prove the maturation of the field. Figure 1 shows some of ontology-related standards that the World Wide Web Consortium (W3C) has adopted in the past decade. Several ontologies and vocabularies have become W3C recommendations: The Time Ontology (OWL-Time) [24]—describing the temporal properties of resources; the Semantic Sensors Network Ontology (SSN) [25]—representing sensors and their observations; the Provenance Ontology (PROV-O) [26]—describing provenance information from different systems; or the RDF Data Cube [27]—enabling publishing of multi-dimensional data on the Web.

Ontology and knowledge representation languages have also evolved as proved by the adoption of new versions of the standards: RDF 1.1 was adopted in February 2014, and introduced identifiers as IRIs, RDF datasets, and new serialization formats, such as RDFa³ and Turtle.⁴ OWL 2.0⁵ was adopted in December

¹<https://data.europeana.eu>

²<http://dati.beniculturali.it/arco/index.php>

³<https://www.w3.org/TR/rdfa-core/>

⁴<https://www.w3.org/TR/turtle/>

⁵<https://www.w3.org/TR/owl2-overview/>

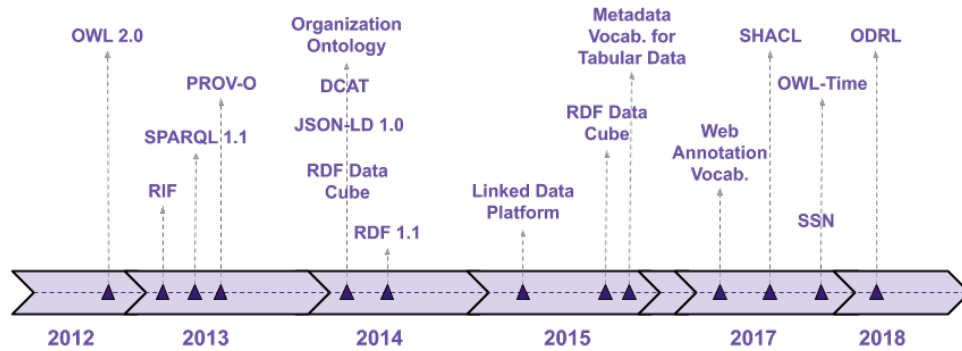


Fig. 1. The timeline of W3C recommendations related to ontologies and vocabularies in the last decade (2010-2019). The years that do not have any recommendations are skipped.

2012, and introduced several new features, such as support for keys and property chains, richer datatypes and data ranges, qualified cardinality restrictions, and enhanced annotation capabilities.

Another notable W3C recommendation adopted in July 2017 is the Shapes Constraints Language (SHACL)⁶ that provides a mechanism for validating constraints against RDF graphs, a feature that was sorely lacking from the current knowledge representation standards. ShEx⁷ is an alternative way of validating RDF and OWL and is backed by an active user community. SHACL and ShEx are considered by some as a simpler knowledge representation languages that might provide an alternative to the more complex OWL representation.

3. Large-Scale Community-Driven Creation of Knowledge

Another area of substantial growth in the last decade is the development of community-authored ontologies and knowledge bases. One of the most visible community-driven projects in the Semantic Web is DBpedia [28]. DBpedia is a crowd-sourced community effort to extract structured, multilingual content from the Wikipedia infoboxes. The English version of the DBpedia knowledge base describes more than 4.5 million things to date.⁸ The DBpedia ontology is rather small—it currently contains 685 classes and

2,795 properties⁹— and it is manually built in a wiki based on the most commonly used information in the infoboxes. To extract the information from the Wikipedia infoboxes, the DBpedia community manually creates the mappings between the infoboxes and the ontology in the DBpedia Mappings Wiki.¹⁰ DBpedia is widely used both within and outside the Semantic Web research community, with many applications and tools being built around it. Most importantly, DBpedia has become a central hub within the web of Open Linked Data, as many RDF data sets link back to DBpedia entities. The Linked Open Data initiative would not have been successful without DBpedia.

A project with a similar goal, but significantly different approach is YAGO [29]. YAGO builds a large-scale ontology also from Wikipedia infoboxes. It uses the Wikipedia categories to find a type for each entity, which is then mapped into the WordNet taxonomy [30]. In this way, YAGO creates a high-quality taxonomy which is used not only in the YAGO-driven applications, but also to perform consistency checks on the automatically extracted information. The YAGO2 ontology [31] extends the YAGO model with spatial and temporal dimensions. The spatial information is obtained by integration with GeoNames,¹¹ a geographical database covering all countries and which contains over eleven million placenames. YAGO2 contains 447 million facts about 9.8 million entities and has a high accuracy of 95% for the facts stored in YAGO. YAGO3 extends the YAGO content with information from the

⁶<https://www.w3.org/TR/shacl/>

⁷<http://shex.io/>

⁸<https://wiki.dbpedia.org/about/facts-figures>

⁹<https://wiki.dbpedia.org/services-resources/ontology>

¹⁰<http://mappings.dbpedia.org>

¹¹<http://www.geonames.org/>

Wikipedias in multiple languages which is fused to create a coherent knowledgebase.

Another notable of large-scale, community-driven knowledge base is Wikidata,¹²—a free and open knowledge base that serves as the central storage of structured data for several Wikimedia projects, including Wikipedia [32]. The Wikidata project started in October 2012. Initially, it was conceived as a central place for storing inter-languages links between Wikipedia articles about the same topic in different languages, and nowadays, Wikidata provides the structured data for almost 60% of Wikipedia pages.¹³ Wikidata's data model is centered around items with unique identifiers that contain statements—basically, key-value pairs—which can be qualified (e.g., with provenance information). One distinguishing feature of Wikidata is its collaborative authoring model: Both humans and programmable bots can contribute content, with a majority of the edits (about 90%) coming from bots. The project is highly active, containing more than 63 million entities, and over 900 million edits as of April 2019.

Another high-impact project for creating vocabularies for structured data to be used on Web content is Schema.org.¹⁴ Started in 2011 by Google, Microsoft, Yahoo and Yandex, the Schema.org vocabularies enable Web content creators to add structured metadata to their Web pages, so that search engines can better understand the content of the page. The Schema.org vocabularies are developed by an open community process using W3C mailing lists and GitHub.¹⁵ Schema.org also offers an extension mechanism that communities have used to create domain-specific vocabularies, for example, for bibliographic or auto extensions. As of April 2019, these extensions are folded back into the main Schema.org vocabulary.¹⁶

Certainly, knowledge graphs (KG) are one of the leading topics of the last decade. Even though researchers have built knowledge networks before, the phrase “knowledge graph” started catching on once Google announced their Google KG in May 2012. Since then, we have seen a flourishing of KGs. Indeed, most large companies, including Amazon, Netflix, Pinterest, LinkedIn, Microsoft, Uber, NASA, IBM, and Alibaba are developing their own KGs. Gartner

also identified knowledge graphs as an emerging technology trend in their 2018 technology report [33]. Even with this high adoption, there is no single widely adopted definition of a KG. A common denominator is that KGs contain entities that are inter-related, and are usually at the data level. The level of formality varies a lot: While some use RDF and OWL and a schema, others are schema-less and use property graphs.¹⁷ Some KGs are built bottom-up using Machine Learning (ML) and Natural Language Processing (NLP) techniques, while others are built top-down. Their uses ranges widely from intelligent search, to analytics, cataloging, data integration, and more.

Community curation is not only used in building large-scale knowledge bases like the ones presented above, but also for building ontologies in different domains. For example, the Gene Ontology (GO) is created by a community-driven workflow, in which community members suggest new entities, or changes to current entities using the Gene Ontology github issue tracker.¹⁸ Only a handful of editors actually edit the GO ontology file based on the community feedback. The GO changes rapidly with nightly releases. Similarly, the ICD-11 ontology project solicits community feedback in the form of comments, structured proposals, or translations through its public Web platform.¹⁹ The proposed changes are discussed by the ICD-11 editorial committee, and if approved, they are implemented in a Web-based collaborative ontology development environment, called iCAT [34]. The ICD-11 project also employs another level of scientific review by sending a PDF extracted from parts of the ontology to domain experts in different medical specialties. Even though several large-scale ontology projects use collaboration processes and involve the larger community in the authoring process, there are no two collaboration workflows that are the same. For more information, we refer the reader to the review by Simperl and Luczak-Rösch [35] of the current collaborative ontology engineering practices.

4. Ontology Engineering

Ontology engineering did not change significantly in the last decade. Even though there has been progress

¹²<https://www.wikidata.org>

¹³http://wdm.wmflabs.org/WD_percentUsageDashboard/

¹⁴<https://schema.org/>

¹⁵<https://www.w3.org/community/schemaorg/>

¹⁶<https://schema.org/docs/extension.html>

¹⁷A property graph is a graph for which the edges are labeled, and both vertices and edges can have any number of key/value properties associated with them.

¹⁸<https://github.com/geneontology/go-ontology/issues>

¹⁹<https://icd.who.int/dev11/l-m/en>

in specific areas, which we will briefly discuss in this section, the work on new ontology engineering methodologies did not seem to progress much. Even to date, the most cited ontology engineering method, according to Google Scholar, is the Ontology 101 guide by Noy and McGuinness [36] from 2001.

The NeOn project (2006-2010) produced the most comprehensive methodology for building networked ontologies [37]. The NeOn methodology describes a set of nine scenarios for building ontologies focusing on reuse of ontological and non-ontological resources, merging, re-engineering, and also accounting for collaboration. In addition, the methodology also publishes a Glossary of Processes and Activities to support collaboration, and methodological guidelines for different processes and activities involved in ontology engineering. Even though the NeOn methodology had modest adoption, the work in the NeOn project produced important research that advanced the field.

In the last decade, researchers have developed other ontology engineering methodologies that have been deployed in specific projects, but are still yet to be widely adopted. For example, the UPON Lite methodology [38] supports the rapid prototyping of trial ontologies, while trying to enhance the role of domain experts and minimize the need for ontology experts. The methodology uses a socially-oriented approach and familiar tools, such as spreadsheets, to make the engineering process more accessible to domain experts. The eXtreme Design (XD) methodology [39] uses an agile approach to ontology engineering that is focused on the reuse of ontology design patterns. The methodology is inspired by the principles of extreme programming and uses a divide-and-conquer approach. XD is iterative and incremental, and tries to address one modeling issue at the time. The modeling issue, defined by a set of competency questions, is mapped to one or more ODPs, which are then integrated into the ontology, and tested using unit tests.

The Gene Ontology (GO)—arguably, the most visible and successful ontology project—has generated several ontology engineering methods and tools that are generic, reusable, and that have already been validated in several large-scale ontology development projects [40, 41]. The OBO Foundry defines many of the principles²⁰ by which OBO Foundry ontologies, including the GO, should abide, such as versioning, naming conventions, defining relations, lo-

cus of authority, documentation, collaboration process, orthogonality of ontologies, and reuse [42]. The OBO Foundry is a good example of how a community can develop ontologies for a specific domain. Motivated by the need to manage ontologies that are becoming more modular and inter-dependent, the GO project developed a continuous integration process using Jenkins and Hudson for building ontologies that became a model for the development of other ontologies [43]. ROBOT²¹ is a generic command-line tool and Java library for performing common ontology tasks, such as, computing differences between ontology versions, merging, extracting ontology modules, reasoning, explanation, materializing inferences, etc. The commands in ROBOT can be chained together to create a powerful, repeatable workflow. Another generic tool that was developed as part of the GO project is TermGenie [44], a Web-based class submission form that can generate new classes, once the submission passes a suite of logical, lexical and structural checks. TermGenie is generic and customizable and has been deployed in the development of several biomedical ontologies.

The adoption of ontologies into mainstream is also proven by the recent publication of several books focusing on ontology engineering, such as, “Demystifying OWL for the Enterprise” in 2018 by Uschold [45], the “Ontology Engineering” in 2019 by Kendall and McGuinness [46], and the “An Introduction to Ontology Engineering” in 2018 by Keet [47].

4.1. Patterns, Templating, and Automation

As ontology engineering became more broadly used, knowledge engineers needed ways to optimize and accelerate parts of the ontology development process. One of the approaches was employing ontology design patterns—small, modular, and reusable solutions to recurrent modeling problems—and templates based on these patterns or other representation regularities in the ontology. Another approach was to use automation, such as bulk imports, or scripts to accelerate ontology population.

The initial work on ontology design patterns (ODP) dates back to 2005 [48]. The research on ODPs has only intensified in the last decade. One of the staples of this area is the ODP repository (<http://www.ontologydesignpatterns.org>), which was developed as

²⁰<http://www.obofoundry.org/principles/fp-000-summary.html>

²¹<http://robot.obolibrary.org/>

part of the NeOn project and it is still actively used. The Workshop on Ontology Design Patterns (WOP) that attracts researchers working on ODPs, as well as users trying to apply them, is already in its 10th edition.²² In their book, “Ontology engineering with ontology design patterns: Foundations and applications”, Hitzler and colleagues [49] provide a current assessment of the research and application of ontology patterns. Some of the new work on ODPs include the definition of a language for the representation of ontology patterns and of their relationships [50].

Several biomedical projects adopted the Dead Simple OWL Design Patterns (DOS-DPs)²³—a lightweight, YAML²⁴-based syntax for specifying design patterns [51]. DOS-DPs support the generation of OWL axioms and user-facing documentation using a simple format that can be parsed using out-of-the-box parsers. With DOS-DP, users can quickly generate new classes, or change existing ones when a design pattern changes.

Other mechanisms for specifying patterns and generating axioms are the Ontology PreProcessing Language (OPPL) [52] and the Tawny OWL [53]. OPPL is a macro language based on the Manchester OWL Syntax [54] that contains instructions for adding or removing entities and axioms to an OWL ontology. Tawny OWL, which is built in Clojure²⁵ and backed by the OWL API [55], provides a programmatic way to build ontologies. Tawny OWL allows ontology engineers to use a wide range of tools available for software development, including versioning, distributed development, building, testing and continuous integration.

Another approach that adopts widely-used technologies from software engineering to ontology development is OntoMaven [56]. OntoMaven adapts the Maven development process to ontology engineering in distributed ontology repositories. It supports the modular reuse of ontologies, versioning, the life cycle and dependency management.

4.2. Better Tooling Is Available

The tooling for building ontologies has also evolved considerably in the last decade. The open-source Protégé ontology editor [57] has grown its active large community to more than 360,000 registered users.

WebProtégé [58] is a Web-based editor for OWL 2.0 with a simplified user interface [59] that supports collaboration. WebProtégé also supports tagging, multilinguality, querying and visualization. The Stanford-hosted WebProtégé server (<https://webprotege.stanford.edu>) hosts more than 60,000 ontology projects that users have created or uploaded to the server.

The OnToology [60]—an open-source project that automates part of the collaborative ontology development process—will generate different types of resources for a GitHub ontology, such as documentation using Widoco [61]; class and taxonomy diagrams using the AR2DTool²⁶; and an evaluation report for common pitfalls using the OOPS! framework [62]. VocBench [63] is another open-source Web-based SKOS editor that focuses on collaboration. There are also other active ontology engineering tools, of which we would like to mention: the Live OWL Documentation Environment (LODE) [64]—an OWL ontology documentation tool that generates documentation in human-readable HTML format, and the WebVOWL [65]—a Web-based tool for visualizing ontologies using a force-directed graph layout, and based on the Visual Notation for OWL Ontologies (VOWL) [66].

Commercial ontology engineering tools have also proliferated and gained wide adoption in the last decade. Some of the commercial offerings include the TopQuadrant’s tool suite²⁷ for vocabulary and metadata management, the PoolParty Semantic Suite,²⁸ or Mondeca’s Intelligent Topic Manager (ITM),²⁹ just to name a few. Gra.fo³⁰ is the most recent addition of commercial ontology tools that was launched in late 2018. Gra.fo is a visual, collaborative, and real-time ontology and knowledge graph schema editor that supports both OWL/RDF and property graphs. Several other commercial ontology tools have morphed recently into knowledge graph solutions.

5. Challenges in Ontology Engineering

Semantic Web languages, and especially OWL, have a steep learning curve [67, 68] and require a change of perspective, especially for people coming

²²<http://ontologydesignpatterns.org/wiki/WOP:2019>

²³https://github.com/INCATools/dead_simple_owl_design_patterns

²⁴<https://yaml.org/>

²⁵<https://clojure.org/>

²⁶<https://github.com/idafens/ar2dtool>

²⁷<https://www.topquadrant.com>

²⁸<https://www.poolparty.biz/>

²⁹<https://mondeca.com/itm/>

³⁰<https://gra.fo>

from software engineering, object-oriented programming, or relational database backgrounds. Newcomers are faced not only with the daunting task of creating a new type of model for their problem or domain, but also trying to find the right tool—ontology editor, visualization, reasoner—and workflow/development process, while the resources for making an informed decision are very scarce.

Many newcomers to semantic technologies start with simpler modeling languages, such as SKOS, for defining their vocabularies, but then struggle to upgrade their model into a more expressive language, such as OWL, as there is no straightforward path between the two languages [69]. Similarly, many application developers start with other types of representations—UML diagrams, mind maps, XML Schemas, spreadsheets—which then need to be converted into OWL. Existing conversion algorithms provide mostly structural transformations based on apriori-defined mapping rules. The conversion results often need to be further processed manually to try to capture the intended semantics of the data source. Although the bootstrapping of ontologies represents a crucial aspect of the ontology development process, and is also the first encounter of newcomers to ontologies—often a make or break issue—it is not as well researched as other related areas, and has less methodological and tool support. Simplifying and better supporting the initial phases of ontology development would encourage a wider adoption of ontologies.

The inherent complexity of OWL also makes it challenging to build developer-friendly APIs for accessing and handling OWL ontologies. The most comprehensive Java API for OWL 2.0 ontologies, the OWL-API [55], requires good knowledge of the OWL specification and it can be intimidating for developers. At the same time, developer-friendly approaches, such as JSON-LD³¹—a lightweight syntax to serialize Linked Data in JSON, or GraphQL³²—a highly popular data query and manipulation language—if used properly, could bolster the adoption of ontologies. There are already several academic and commercial approaches that use GraphQL to query RDF graphs. Even so, more research is needed to tackle the challenges of bridging GraphQL, RDF and SPARQL: GraphQL does not have a notion of semantics, nor global identifiers, similar to the URIs in RDF, while the query results in GraphQL represent trees rather than full graphs [70].

³¹<https://json-ld.org/>

³²<https://graphql.org/>

The adoption of ontologies is also hindered by competing approaches that are simpler to use than OWL. Microdata and microformats were much wider found in the 2013 Common Crawl dataset than RDFa [71]. Schema.org chose to use an extension of RDFS Schema as its data model,³³ rather than OWL, and introduced `domainIncludes` and `rangeIncludes` (as alternatives for the `rdfs:domain` and `rdfs:range`) for pragmatic reasons, so that it is easier to encode multiple values for the domain and ranges of properties. ShACL and ShEx, while being developed to validate graph-based data against a set of conditions, is seen by some in the community as a simpler knowledge representation language that might one day replace OWL. Many commercial companies are choosing labeled property graphs to encode their knowledge graphs rather than RDF triplestores as the property graphs are perceived to be simpler to understand, and usually offer better query performance.

One key aspect of ontology engineering is reuse, both of ontologies and of parts of ontologies. Even though, several ontology repositories exist [72]—BioPortal [73] for biomedicine, AgroPortal [74] for agriculture, or general-purpose repositories, such as the Linked Open Vocabularies (LOV) [75] and Ontohub [76]—finding the right ontology for a particular task is still difficult. The testament to this challenge are the countless postings on the Semantic Web mailing lists from interested users trying to find an ontology for a particular domain or task. A common scenario is finding an ontology that can be used to annotate a corpus of text. The BioPortal Ontology Recommender³⁴ [77] uses several criteria to make ontology recommendations for the biomedical domain, however a general-purpose recommender is not available. One of the bigger research challenges in finding the “right” ontology is coming up with a set of metrics that can measure how suitable an ontology is for a task or a domain—a topic that has been one of the main subjects of ontology evaluation for a long time, but it is still not solved. While a review of current ontology evaluation methods is not covered here, the reader can refer to the review by Sabou and Fernandez for further details [78], and to a more recent description of the state of the art in ontology evaluation by Poveda Villalon [79].

Reusing parts of an ontology—single entities, subtrees, or sets of axioms—is often necessary during on-

³³<https://schema.org/docs/datamodel.html>

³⁴<https://bioportal.bioontology.org/recommender>

1 tology development, however enacting reuse in prac-
 2 tice is difficult. Several published studies have shown
 3 that the level of ontological reuse is low [80–82]. For
 4 example, in a recent study, Kamdar et. al [81] show
 5 that the term reuse is less than 9% in biomedical on-
 6 tologies, even though the term overlap is between 25–
 7 31%, with most ontologies reusing fewer than 5% of
 8 their terms from a small set of popular ontologies. The
 9 challenges related to reuse come for a wide range:
 10 from finding the ontology to reuse, to extracting the
 11 subset to reuse (although several module-extraction al-
 12 gorithms exist [83]), to maintaining the extracted sub-
 13 set as the source ontology evolves.

14 Due to space limitations, the challenges described in
 15 this section do not represent a comprehensive listing of
 16 challenges in ontology engineering, but rather repre-
 17 sent issues that have been encountered all too often by
 18 the author and her collaborators. The interested reader
 19 can learn more about the current state and challenges
 20 in using ontology design pattern from Blomqvist et.
 21 al [84], about the state of ontology evolution from
 22 Zablith et. al [85], and about ontology matching from
 23 Otero-Cerdeira et. al [86]. Shaviko et. al [87] cover the
 24 state and challenges of ontology learning, while a more
 25 introductory and comprehensive view can be found in
 26 the book “Perspectives on Ontology Learning” edited
 27 by Lehmann and Völker [88].

30 6. Opportunities for Future Research

32 Even though the research topics in Semantic Web
 33 have evolved in the last decade, ontologies were a con-
 34 stant presence, as shown in Figure 2. The word clouds
 35 in Figure 2 were generated from the titles of the ac-
 36 cepted papers at the International Semantic Web Con-
 37 ference (ISWC) in three different years of the last
 38 decade, and arguably represent a fairly accurate im-
 39 age of the research topics at the time. An investigation
 40 of the topics in the call for papers for larger Sema-
 41 ntic Web conferences also confirms that ontologies and
 42 their engineering have always been present in the last
 43 decade. On one hand, this constancy points to how im-
 44 portant ontologies are for the Semantic Web, and on
 45 the other hand, it is a sign that ontology engineering is
 46 still an active research topic that needs to evolve sig-
 47 nificantly in the next decade.


48 One of main show stoppers for a wider adoption of
 49 ontologies is their steep learning curve and the com-
 50 plexity of the languages and tooling, as we discussed
 51 in the previous section. There is a stringent need to

1 simplify, if not the actual Semantic Web languages
 2 and standards, but at least the presentation and inter-
 3 action of the users with the Semantic Web languages
 4 and tooling. The Semantic Web community can ben-
 5 efit significantly from a tighter collaboration with the
 6 Human Computer Interaction (HCI) communities and
 7 from applying HCI techniques in the design of their
 8 tooling. Two encouraging, albeit rare examples, are the
 9 paper by Fu et. al [89] that investigates ontology visu-
 10 alization techniques in the context of class mappings
 11 and the paper by Vigo et. al [68] that provides design
 12 guidelines for ontology editors. Displaying class hier-
 13 archies, which are so central to ontology development
 14 but are still cumbersome to use, is another area where
 15 HCI already provides several solutions [90].

16 At the same time, the user interfaces for eliciting the
 17 content of an ontology have not evolved much. Cur-
 18 rent ontology editors, including Protégé, offer an in-
 19 timidating view of all features that OWL offers. It is
 20 no wonder that newcomers are scared away. We need
 21 role-based user interfaces that would enable users with
 22 different expertise to contribute efficiently. These user
 23 interfaces could be automatically generated based on a
 24 user profile, and at the same time the interfaces need
 25 to enforce certain editing rules (e.g., all classes need to
 26 have a `rdfs:label`) that are paramount in any real-
 27 world ontology development project.

28 Another opportunity for simplifying the presenta-
 29 tion and consumption of ontologies is to continue the
 30 research on ontology summarization. Although a few
 31 summarization approaches exist already [91, 92], there
 32 are still several open issues. Most notably it is not clear
 33 what is the best way to evaluate such approaches. Fu-
 34 ture directions in ontology summarization include the
 35 customization of the summarization by allowing the
 36 user to tune the model to generate different summaries
 37 based on different requirements, as well as, using non-
 38 extractive techniques, in which the summary is not us-
 39 ing necessarily terms extracted from the ontology [91].
 40 Effective summarization techniques could have a great
 41 impact in helping disseminate ontologies to other com-
 42 munities, and to help our community better find and
 43 reuse ontologies.

44 Knowledge graphs are gaining in popularity, and
 45 they will likely become even more widespread in the
 46 near future. Knowledge graphs can be seen both as a
 47 challenge and as an opportunity with respect to on-
 48 tologies. Knowledge graphs are now widely adopted
 49 in industry, however many of these KGs do not use
 50 RDF, triplestores or semantics, but rather ML tech-
 51 niques for automatic building and property graphs



ity³⁵ co-located with ISWC 2019, and a similar session which took place at the 2nd US Semantic Technologies Symposium (US2TS 2019)³⁶. Explainable AI is trying to produce techniques that will allow human users to understand how complex AI and ML-based systems reach a decision, i.e., find an explanation of a decision, help us produce more explainable models, and enable us to debug the decisions.³⁷ Ontologies have the potential of playing an important role into making AI systems explainable as they already provide a user's conceptualization of a domain, which could be used as part of the explanation or debugging process. However, exactly how to achieve this is a difficult and open issue. We will need new design patterns (some initial work already exists [96]) and even new methodologies for building ontologies that can support explainable systems. We will need to define the interplay of ontologies with the different AI techniques, such as deep learning methods. There have already been some efforts in this direction, referred to as neural-symbolic integration [97], but it is still a new field of investigation that if successful, will have a high impact on society.

The goal of the paper is to give an overview of how the field of ontology engineering has evolved in the last decade. Due to space limitations, we could only cover some of the main topics in ontology engineering. We hope that the paper can serve as an entry point for a newcomer in the ontology field,

<https://www.darpa.mil/program/explainable-artificial-intelligence>

and as a quick reference for the more knowledgeable researchers. As a result of the research and development efforts in the last ten years, ontologies are now adopted in wide range of domains, from biomedicine to engineering and finance. The infrastructures for storing, finding and building ontologies have also evolved significantly. Several standards pertaining to ontology engineering have been adopted in the last decade, and highly-visible efforts to build large-scale ontologies and knowledge bases are well underway. Even though the ontology engineering field still faces several challenges—many of them longstanding—we have also identified many opportunities for future research and development, and exciting new opportunities from synergies with other domains that can drive the ontology engineering field even further.

8. Acknowledgements

The author would like to thank the reviewers for providing helpful feedback, and to Jennifer Vendetti and Csongor Nyulas for proof-reading the paper.

References

- [1] T.R. Gruber, A translation approach to portable ontology specifications., *Knowledge Acquisition* **5**(2) (1993), 199–220.
- [2] R. Stevens, A. Rector and D. Hull, What is an ontology?, *Ontogenesis* (2010). <http://ontogenesis.knowledgeblog.org/66>.
- [3] M. Ashburner, C.A. Ball, J.A. Blake, D. Botstein, H. Butler, J.M. Cherry, A.P. Davis, K. Dolinski, S.S. Dwight, J.T. Eppig et al., Gene ontology: tool for the unification of biology, *Nature genetics* **25**(1) (2000), 25.
- [4] P. Khatri, M. Sirota and A.J. Butte, Ten years of pathway analysis: current approaches and outstanding challenges, *PLoS computational biology* **8**(2) (2012), e1002375.
- [5] R. Hoehndorf, P.N. Schofield and G.V. Gkoutos, The role of ontologies in biological and biomedical research: a functional perspective, *Briefings in bioinformatics* **16**(6) (2015), 1069–1080.
- [6] M. Ivanović and Z. Budimac, An overview of ontologies and data resources in medical domains, *Expert Systems with Applications* **41**(11) (2014), 5158–5166.
- [7] W.H. Organization, International Classification of Diseases (ICD) Information Sheet, 2019, [Online; accessed 05-May-2019]. <https://www.who.int/classifications/icd/factsheet/en/>.
- [8] T. Tudorache, C. Nyulas, N.F. Noy and M.A. Musen, Using Semantic Web in ICD-11: Three Years Down the Road, in: *The Semantic Web - ISWC 2013 - 12th International Semantic Web Conference, Sydney, NSW, Australia, October 21-25, 2013, Proceedings, Part II*, 2013, pp. 195–211. doi:10.1007/978-3-642-41338-4_13.
- [9] M. Bennett, The financial industry business ontology: Best practice for big data, *Journal of Banking Regulation* **14**(3–4) (2013), 255–268.
- [10] W. Borst, Construction of engineering ontology, *Ph. D. Dissertation, Institute for Telematica and Information Technology, University of Twente, Enschede, the Netherlands* (1997).
- [11] T.R. Gruber and G.R. Olsen, An ontology for engineering mathematics, in: *Principles of Knowledge Representation and Reasoning*, Elsevier, 1994, pp. 258–269.
- [12] S. Feldmann, S. Rösch, C. Legat and B. Vogel-Heuser, Keeping requirements and test cases consistent: towards an ontology-based approach, in: *2014 12th IEEE International Conference on Industrial Informatics (INDIN)*, IEEE, 2014, pp. 726–732.
- [13] W. Terkaj and M. Urgo, Ontology-based modeling of production systems for design and performance evaluation, in: *2014 12th IEEE International Conference on Industrial Informatics (INDIN)*, IEEE, 2014, pp. 748–753.
- [14] O. Kovalenko, E. Serral, M. Sabou, F.J. Ekaputra, D. Winkler and S. Biffl, Automating cross-disciplinary defect detection in multi-disciplinary engineering environments, in: *International Conference on Knowledge Engineering and Knowledge Management*, Springer, 2014, pp. 238–249.
- [15] S. Biffl and M. Sabou, *Semantic web technologies for intelligent engineering applications*, Springer, 2016.
- [16] E. Hyvönen, *Publishing and Using Cultural Heritage Linked Data on the Semantic Web*, Synthesis Lectures on the Semantic Web, Morgan & Claypool Publishers, 2012. doi:10.2200/S00452ED1V01Y201210WBE003.
- [17] A. Isaac and B. Haslhofer, Europeana linked open data—data. europeana.eu, *Semantic Web* **4**(3) (2013), 291–297.
- [18] S. Peroni and D. Shotton, The SPAR ontologies, in: *International Semantic Web Conference*, Springer, 2018, pp. 119–136.
- [19] V.A. Carriero, A. Gangemi, M.L. Mancinelli, L. Marinucci, A.G. Nuzzolese, V. Presutti and C. Veninata, ArCo: the Italian Cultural Heritage Knowledge Graph, in: *The Semantic Web - ISWC 2019 - 18th International Semantic Web Conference, Auckland, New Zealand, October 26-30, 2019, Proceedings*, 2019.
- [20] International Organization for Standardization (ISO), ISO 21127:2014 Information and documentation — A reference ontology for the interchange of cultural heritage information, 2014, [Online; accessed 08-Oct-2019]. <https://www.iso.org/standard/57832.html>.
- [21] C.M. de Oliveira Rodrigues, F.L.G. de Freitas, E.F.S. Barreiros, R.R. de Azevedo and A.T. de Almeida Filho, Legal ontologies over time: a systematic mapping study, *Expert Systems with Applications* **130** (2019), 12–30.
- [22] International Organization for Standardization (ISO), ISO/TS 19150-1:2012 Geographic information — Ontology — Part 1: Framework, 2012, [Online; accessed 08-Oct-2019]. <https://www.iso.org/standard/57465.html>.
- [23] International Organization for Standardization (ISO), ISO 19150-2:2015 Geographic information — Ontology — Part 2: Rules for developing ontologies in the Web Ontology Language (OWL), 2015, [Online; accessed 08-Oct-2019]. <https://www.iso.org/standard/57466.html>.
- [24] W.W.W. Consortium, Time Ontology in OWL, W3C Recommendation, 2017, [Online; accessed 05-May-2019]. <https://www.w3.org/TR/owl-time/>.

- [25] W.W.W. Consortium, Semantic Sensor Network Ontology, W3C Recommendation, 2017, [Online; accessed 05-May-2019]. <https://www.w3.org/TR/vocab-ssn/>.
- [26] W.W.W. Consortium, PROV-O: The PROV Ontology, W3C Recommendation, 2013, [Online; accessed 05-May-2019]. <https://www.w3.org/TR/prov-o/>.
- [27] W.W.W. Consortium, The RDF Data Cube Vocabulary, W3C Recommendation, 2014, [Online; accessed 05-May-2019]. <https://www.w3.org/TR/2014/REC-vocab-data-cube-20140116/>.
- [28] J. Lehmann, R. Isele, M. Jakob, A. Jentzsch, D. Kontokostas, P.N. Mendes, S. Hellmann, M. Morsey, P. Van Kleef, S. Auer et al., DBpedia—a large-scale, multilingual knowledge base extracted from Wikipedia, *Semantic Web* **6**(2) (2015), 167–195.
- [29] F.M. Suchanek, G. Kasneci and G. Weikum, Yago: a core of semantic knowledge, in: *Proceedings of the 16th international conference on World Wide Web*, ACM, 2007, pp. 697–706.
- [30] G.A. Miller, WordNet: a lexical database for English, *Communications of the ACM* **38**(11) (1995), 39–41.
- [31] J. Hoffart, F.M. Suchanek, K. Berberich and G. Weikum, YAGO2: A spatially and temporally enhanced knowledge base from Wikipedia, *Artificial Intelligence* **194** (2013), 28–61.
- [32] D. Vrandečić and M. Krötzsch, Wikidata: a free collaborative knowledgebase, *Commun. ACM* **57**(10) (2014), 78–85. doi:10.1145/2629489.
- [33] Gartner, Gartner Identifies Five Emerging Technology Trends That Will Blur the Lines Between Human and Machine, 2018, [Online; accessed 05-May-2019]. <https://tinyurl.com/yxbxdk3>.
- [34] T. Tudorache, S.M. Falconer, N.F. Noy, C. Nyulas, T.B. Üstün, M.D. Storey and M.A. Musen, Ontology Development for the Masses: Creating ICD-11 in WebProtégé, in: *Knowledge Engineering and Management by the Masses - 17th International Conference, EKAW 2010, Lisbon, Portugal, October 11-15, 2010. Proceedings*, 2010, pp. 74–89. doi:10.1007/978-3-642-16438-5_6.
- [35] E. Simperl and M. Luczak-Rösch, Collaborative ontology engineering: a survey, *Knowledge Eng. Review* **29**(1) (2014), 101–131. doi:10.1017/S0269888913000192.
- [36] N.F. Noy, D.L. McGuinness et al., Ontology development 101: A guide to creating your first ontology, Stanford knowledge systems laboratory technical report KSL-01-05, 2001.
- [37] M.C. Suárez-Figueroa, A. Gómez-Pérez and M. Fernández-López, The NeOn methodology for ontology engineering, in: *Ontology engineering in a networked world*, Springer, 2012, pp. 9–34.
- [38] A. De Nicola and M. Missikoff, A lightweight methodology for rapid ontology engineering, *Communications of the ACM* **59**(3) (2016), 79–86.
- [39] E. Blomqvist, V. Presutti, E. Daga and A. Gangemi, Experimenting with eXtreme design, in: *International Conference on Knowledge Engineering and Knowledge Management*, Springer, 2010, pp. 120–134.
- [40] G.O. Consortium, Gene ontology consortium: going forward, *Nucleic acids research* **43**(D1) (2014), D1049–D1056.
- [41] C. Mungall, H. Dietze and D. Osumi-Sutherland, Use of OWL within the Gene Ontology, in: *Proceedings of the 11th International Workshop on OWL: Experiences and Directions (OWLED 2014) co-located with 13th International Semantic Web Conference on (ISWC 2014)*, Riva del Garda, Italy, October 17-18, 2014., 2014, pp. 25–36. http://ceur-ws.org/Vol-1265/owled2014_submission_5.pdf.
- [42] B. Smith, M. Ashburner, C. Rosse, J. Bard, W. Bug, W. Ceusters, L.J. Goldberg, K. Eilbeck, A. Ireland, C.J. Mungall et al., The OBO Foundry: coordinated evolution of ontologies to support biomedical data integration, *Nature biotechnology* **25**(11) (2007), 1251.
- [43] C.J. Mungall, H. Dietze, S.J. Carbon, A. Ireland, S. Bauer and S. Lewis, Continuous integration of open biological ontology libraries, *Bio-Ontologies 2012* (2012).
- [44] H. Dietze, T.Z. Berardini, R.E. Foulger, D.P. Hill, J. Lomax, D. Osumi-Sutherland, P. Roncaglia and C.J. Mungall, TermGenie—a web-application for pattern-based ontology class generation, *Journal of biomedical semantics* **5**(1) (2014), 48.
- [45] M. Uschold, Demystifying OWL for the Enterprise, *Synthesis Lectures on Semantic Web: Theory and Technology* **8**(1) (2018), i–237.
- [46] E.F. Kendall and D. McGuinness, Ontology Engineering, *Synthesis Lectures on Semantic Web: Theory and Technology* **18**(1) (2019), i–102.
- [47] M. Keet, An Introduction to Ontology Engineering, 2018, Accessed online on May 25, 2019. <http://www.meteck.org/teaching/OEbook/>.
- [48] A. Gangemi, Ontology design patterns for semantic web content, in: *International semantic web conference*, Springer, 2005, pp. 262–276.
- [49] P. Hitzler, A. Gangemi, K. Janowicz, A.A. Krisnadhi and V. Presutti, *Ontology engineering with ontology design patterns: Foundations and applications*, Vol. 25, IOS Press, 2016.
- [50] P. Hitzler, A. Gangemi, K. Janowicz, A.A. Krisnadhi and V. Presutti, Towards a Simple but Useful Ontology Design Pattern Representation Language., in: *WOP@ ISWC*, 2017.
- [51] D. Osumi-Sutherland, M. Courtot, J.P. Balhoff and C. Mungall, Dead simple OWL design patterns, *Journal of biomedical semantics* **8**(1) (2017), 18.
- [52] M. Egana, R. Stevens and E. Antezana, Transforming the Axiomatisation of Ontologies: The Ontology Pre-Processor Language., *OWLED (Spring)* **496** (2008).
- [53] P. Lord, The Semantic Web takes Wing: Programming Ontologies with Tawny-OWL, in: *Proceedings of the 10th International Workshop on OWL: Experiences and Directions (OWLED 2013) co-located with 10th Extended Semantic Web Conference (ESWC 2013)*, Montpellier, France, May 26-27, 2013., 2013. http://ceur-ws.org/Vol-1080/owled2013_16.pdf.
- [54] M. Horridge and P.F. Patel-Schneider, OWL 2 web ontology language Manchester syntax, *W3C Working Group Note* (2009).
- [55] M. Horridge and S. Bechhofer, The OWL api: A Java API for OWL ontologies, *Semantic Web* **2**(1) (2011), 11–21.
- [56] A. Paschke and R. Schäfermeier, OntoMaven-Maven-Based Ontology Development and Management of Distributed Ontology Repositories, in: *Synergies Between Knowledge Engineering and Software Engineering*, Springer, 2018, pp. 251–273.
- [57] M.A. Musen, The protégé project: a look back and a look forward, *AI Matters* **1**(4) (2015), 4–12. doi:10.1145/2757001.2757003.
- [58] M. Horridge, R.S. Gonçalves, C.I. Nyulas, T. Tudorache and M.A. Musen, WebProtégé: A Cloud-Based Ontology Editor, in: *Companion of The 2019 World Wide Web Conference, WWW 2019, San Francisco, CA, USA, May 13-17, 2019.*, 2019, pp. 686–689. doi:10.1145/3308560.3317707.

- [59] M. Horridge, T. Tudorache, J. Vendetti, C. Nyulas, M.A. Musen and N.F. Noy, Simplified OWL Ontology Editing for the Web: Is WebProtégé Enough?, in: *The Semantic Web - ISWC 2013 - 12th International Semantic Web Conference, Sydney, NSW, Australia, October 21-25, 2013, Proceedings, Part 1*, 2013, pp. 200–215. doi:10.1007/978-3-642-41335-3_13.
- [60] A. Alobaid, D. Garijo, M. Poveda-Villalón, I. Santana-Perez, A. Fernández-Izquierdo and O. Corcho, Automating ontology engineering support activities with OnToology, *Journal of Web Semantics* (2018).
- [61] D. Garijo, WIDOCO: a wizard for documenting ontologies, in: *International Semantic Web Conference*, Springer, 2017, pp. 94–102.
- [62] M. Poveda-Villalón, A. Gómez-Pérez and M.C. Suárez-Figueroa, Oops!(ontology pitfall scanner!): An on-line tool for ontology evaluation, *International Journal on Semantic Web and Information Systems (IJSWIS)* **10**(2) (2014), 7–34.
- [63] A. Stellato, S. Rajbhandari, A. Turbati, M. Fiorelli, C. Caracciolo, T. Lorenzetti, J. Keizer and M.T. Pazeniza, VocBench: a web application for collaborative development of multilingual thesauri, in: *European Semantic Web Conference*, Springer, 2015, pp. 38–53.
- [64] S. Peroni, D. Shotton and F. Vitali, Tools for the automatic generation of ontology documentation: A task-based evaluation, *International Journal on Semantic Web and Information Systems (IJSWIS)* **9**(1) (2013), 21–44.
- [65] S. Lohmann, V. Link, E. Marbach and S. Negru, WebVOWL: Web-based visualization of ontologies, in: *International Conference on Knowledge Engineering and Knowledge Management*, Springer, 2014, pp. 154–158.
- [66] S. Lohmann, S. Negru, F. Haag and T. Ertl, Visualizing ontologies with VOWL, *Semantic Web* **7**(4) (2016), 399–419.
- [67] P. Warren, P. Mulholland, T. Collins and E. Motta, The usability of description logics, in: *European Semantic Web Conference*, Springer, 2014, pp. 550–564.
- [68] M. Vigo, S. Bail, C. Jay and R.D. Stevens, Overcoming the pitfalls of ontology authoring: Strategies and implications for tool design, *Int. J. Hum.-Comput. Stud.* **72**(12) (2014), 835–845. doi:10.1016/j.ijhcs.2014.07.005.
- [69] S. Jupp, S. Bechhofer and R. Stevens, SKOS with OWL: Don't be Full-ish!, in: *Proceedings of the Fifth OWLED Workshop on OWL: Experiences and Directions, collocated with the 7th International Semantic Web Conference (ISWC-2008), Karlsruhe, Germany, October 26-27, 2008*, 2008. http://ceur-ws.org/Vol-432/owlled2008eu_submission_22.pdf.
- [70] R. Taelman, M.V. Sande and R. Verborgh, GraphQL-LD: Linked Data Querying with GraphQL, in: *Proceedings of the ISWC 2018 Posters & Demonstrations, Industry and Blue Sky Ideas Tracks co-located with 17th International Semantic Web Conference (ISWC 2018), Monterey, USA, October 8th - to - 12th, 2018*, 2018. <http://ceur-ws.org/Vol-2180/paper-65.pdf>.
- [71] R. Meusel, P. Petrovski and C. Bizer, The webdatacommons microdata, rdfa and microformat dataset series, in: *International Semantic Web Conference*, Springer, 2014, pp. 277–292.
- [72] M. d'Áquin and N.F. Noy, Where to publish and find ontologies? A survey of ontology libraries, *Journal of Web Semantics* **11** (2012), 96–111.
- [73] P.L. Whetzel, N.F. Noy, N.H. Shah, P.R. Alexander, C. Nyulas, T. Tudorache and M.A. Musen, BioPortal: enhanced functionality via new Web services from the National Center for Biomedical Ontology to access and use ontologies in software applications, *Nucleic acids research* **39**(suppl_2) (2011), W541–W545.
- [74] C. Jonquet, A. Toulet, E. Arnaud, S. Aubin, E.D.Y. Kaboré, V. Emonet, J. Graybeal, M. Laporte, M.A. Musen, V. Pesce and P. Larmande, AgroPortal: A vocabulary and ontology repository for agronomy, *Computers and Electronics in Agriculture* **144** (2018), 126–143. doi:10.1016/j.compag.2017.10.012.
- [75] P.-Y. Vandenbussche, G.A. Atemezing, M. Poveda-Villalón and B. Vatant, Linked Open Vocabularies (LOV): a gateway to reusable semantic vocabularies on the Web, *Semantic Web* **8**(3) (2017), 437–452.
- [76] T. Mossakowski, O. Kutz and M. Codescu, Ontohub: A semantic repository for heterogeneous ontologies, in: *IN: PROC. OF DACSăĂŽ14 (2014)*, Citeseer, 2014.
- [77] M.M. Romero, C. Jonquet, M.J. O'Connor, J. Graybeal, A. Pazos and M.A. Musen, NCBO Ontology Recommender 2.0: an enhanced approach for biomedical ontology recommendation, *J. Biomedical Semantics* **8**(1) (2017), 21:1–21:22. doi:10.1186/s13326-017-0128-y.
- [78] M. Sabou and M. Fernandez, Ontology (network) evaluation, in: *Ontology engineering in a networked world*, Springer, 2012, pp. 193–212.
- [79] M.P. Villalon, Ontology evaluation: a pitfall-based approach to ontology diagnosis, PhD thesis, Universidad Politecnica de Madrid, 2016.
- [80] A. Ghazvinian, N.F. Noy and M.A. Musen, How orthogonal are the OBO Foundry ontologies?, in: *Journal of biomedical semantics*, Vol. 2, BioMed Central, 2011, p. S2.
- [81] M.R. Kamdar, T. Tudorache and M.A. Musen, A systematic analysis of term reuse and term overlap across biomedical ontologies, *Semantic web* **8**(6) (2017), 853–871.
- [82] M. Poveda Villalón, M.C. Suárez-Figueroa and A. Gómez-Pérez, The landscape of ontology reuse in linked data (2012).
- [83] M. d'Áquin, Modularizing ontologies, in: *Ontology Engineering in a Networked World*, Springer, 2012, pp. 213–233.
- [84] E. Blomqvist, P. Hitzler, K. Janowicz, A. Krisnadhi, T. Narock and M. Solanki, Considerations regarding Ontology Design Patterns., *Semantic Web* **7**(1) (2016), 1–7.
- [85] F. Zablith, G. Antoniou, M. d'Áquin, G. Flouris, H. Kondylakis, E. Motta, D. Plexousakis and M. Sabou, Ontology evolution: a process-centric survey, *The knowledge engineering review* **30**(1) (2015), 45–75.
- [86] L. Otero-Cerdeira, F.J. Rodríguez-Martínez and A. Gómez-Rodríguez, Ontology matching: A literature review, *Expert Systems with Applications* **42**(2) (2015), 949–971.
- [87] P. Shvaiko and J. Euzenat, Ontology matching: state of the art and future challenges, *IEEE Transactions on knowledge and data engineering* **25**(1) (2011), 158–176.
- [88] J. Lehmann and J. Völker, *Perspectives on ontology learning*, Vol. 18, IOS Press, 2014.
- [89] B. Fu, N.F. Noy and M.-A. Storey, Indented tree or graph? A usability study of ontology visualization techniques in the context of class mapping evaluation, in: *International Semantic Web Conference*, Springer, 2013, pp. 117–134.
- [90] H.-J. Schulz, Trevis. net: A tree visualization reference, *IEEE Computer Graphics and Applications* **31**(6) (2011), 11–15.

- [91] S. Pouriyeh, M. Allahyari, Q. Liu, G. Cheng, H.R. Arabnia, M. Atzori and K. Kochut, Graph-based methods for ontology summarization: A survey, in: *2018 IEEE First International Conference on Artificial Intelligence and Knowledge Engineering (AIKE)*, IEEE, 2018, pp. 85–92.
- [92] N. Li and E. Motta, Evaluations of user-driven ontology summarization, in: *International Conference on Knowledge Engineering and Knowledge Management*, Springer, 2010, pp. 544–553.
- [93] Q. Wang, Z. Mao, B. Wang and L. Guo, Knowledge graph embedding: A survey of approaches and applications, *IEEE Transactions on Knowledge and Data Engineering* **29**(12) (2017), 2724–2743.
- [94] O.M. Holter, E.B. Myklebust, J. Chen and E. Jimenez-Ruiz, Embedding OWL Ontologies with OWL2Vec*, in: *Proceedings of the ISWC 2019 Satellite Tracks (Posters Demonstrations, Industry, and Outrageous Ideas) co-located with 18th International Semantic Web Conference (ISWC 2019)*, Vol. 2456, CEUR-WS, 2019.
- [95] P. Kolyvakis, A. Kalousis and D. Kiritsis, DeepAlignment: Un-supervised ontology matching with refined word vectors, in: *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, 2018, pp. 787–798.
- [96] I. Tiddi, M. d’Aquin and E. Motta, An ontology design pattern to define explanations, in: *Proceedings of the 8th International Conference on Knowledge Capture*, ACM, 2015, p. 3.
- [97] T.R. Besold, A.d. Garcez, S. Bader, H. Bowman, P. Domingos, P. Hitzler, K.-U. Kühnberger, L.C. Lamb, D. Lowd, P.M.V. Lima et al., Neural-symbolic learning and reasoning: A survey and interpretation, *arXiv preprint arXiv:1711.03902* (2017).